

Lab 4 Report

Summary:

For the serial version of the program, I got 4.04 Gigaflops per second, and for the CUDA implementation, I got 499.29 Gigaflops per second. For my CUDA compute structure I used a single block with 1024 threads. Since the operation was easily parallelizable, the amount of work done by the CUDA implementation was over 100 times faster, which makes sense because of the division of work we were able to achieve in our program.